Felix_Week2_Assign_Pima

August 17, 2022

[ ]:
```
#Q 1. Import the necessary libraries and briefly explain the use of each library

# import the important packages
import pandas as pd # library used for data manipulation and analysis
import numpy as np # library used for working with arrays
import matplotlib.pyplot as plt # library for plots and visualisations
import seaborn as sns # library for visualisations

%matplotlib inline

import scipy.stats as stats # this library contains a large number of␣
 ↪probability distributions as well as

from scipy.stats import norm # this library is used for normal distribution
```

[186]:
```
#Q3. Show the last 10 records of the dataset. How many columns are there?
# There are 9 Columns
diabetes = pd.read_csv("diabetes.csv")
diabetes.tail(10)
diabetes.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
 #   Column                    Non-Null Count  Dtype
---  ------                    --------------  -----
 0   Pregnancies               768 non-null    int64
 1   Glucose                   768 non-null    int64
 2   BloodPressure             768 non-null    int64
 3   SkinThickness             768 non-null    int64
 4   Insulin                   768 non-null    int64
 5   BMI                       768 non-null    float64
 6   DiabetesPedigreeFunction  768 non-null    float64
 7   Age                       768 non-null    int64
 8   Outcome                   768 non-null    int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

[162]:
```
#Q4. Show the first 10 records of the dataset
diabetes.head(10)
```

[162]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	6	148	72	35	79	33.600000	0.627	50	1
1	1	85	66	29	79	26.600000	0.351	31	0
2	8	183	64	20	79	23.300000	0.672	32	1
3	1	89	66	23	94	28.100000	0.167	21	0
4	0	137	40	35	168	43.100000	2.288	33	1
5	5	116	74	20	79	25.600000	0.201	30	0
6	3	78	50	32	88	31.000000	0.248	26	1
7	10	115	69	20	79	35.300000	0.134	29	0
8	2	197	70	45	543	30.500000	0.158	53	1
9	8	125	96	20	79	31.992578	0.232	54	1

[178]:
```
# Q5. What do you understand by the dimension of the dataset? Find the␣
 ↪dimension of the `pima` dataframe.
# What I understand by its dimension of the dataset is it is a DataFrame since␣
 ↪being two-dimensional It contains 768 rows
# and 9 colums
diabetes.ndim
```

[178]: 2

[182]:
```
diabetes.shape
```

[182]: (768, 9)

[181]:
```
# Q6. What do you understand by the size of the dataset? Find the size of the␣
 ↪`pima` dataframe.
# since the size of the dataset is 6912, it is composed of 768 rows and 9␣
 ↪columns.
diabetes.size
```

[181]: 6912

[190]:
```
#Q7. What are the data types of all the variables in the data set?
diabetes.dtypes
```

[190]:
```
Pregnancies                   int64
Glucose                       int64
BloodPressure                 int64
SkinThickness                 int64
Insulin                       int64
BMI                         float64
DiabetesPedigreeFunction    float64
Age                           int64
Outcome                       int64
dtype: object
```

[ ]:
```

```

[197]:
```
# Q8 What do you mean by missing values? Are there any missing values in the␣
 ↪`pima` dataframe?
# missing values of cells that are voided of values. There are no missing␣
 ↪values in the dataset.
diabetes.isnull()
```

[197]:
```
     Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin    BMI  \
0          False    False          False          False    False  False
1          False    False          False          False    False  False
2          False    False          False          False    False  False
3          False    False          False          False    False  False
4          False    False          False          False    False  False
..           …        …              …              …        …      …
763        False    False          False          False    False  False
764        False    False          False          False    False  False
765        False    False          False          False    False  False
766        False    False          False          False    False  False
767        False    False          False          False    False  False

     DiabetesPedigreeFunction    Age  Outcome
0                       False  False    False
1                       False  False    False
2                       False  False    False
3                       False  False    False
4                       False  False    False
..                        …      …        …
763                     False  False    False
764                     False  False    False
765                     False  False    False
766                     False  False    False
767                     False  False    False
```

```
[768 rows x 9 columns]
```

[191]:
```
# Q9. What does summary statistics of data represents? Find the summary␣
 ↪statistics for all variables except 'Outcome'
#in the `pima` data? Take one column/variable from the output table and explain␣
 ↪all the statistical measures.

# Summary statistics give an overview or summary of descritive statistics on␣
 ↪the variables of the sataset.
diabetes.describe()
```

[191]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	\
count	768.000000	768.000000	768.000000	768.000000	768.000000	
mean	3.845052	121.675781	72.250000	26.447917	118.270833	
std	3.369578	30.436252	12.117203	9.733872	93.243829	
min	0.000000	44.000000	24.000000	7.000000	14.000000	
25%	1.000000	99.750000	64.000000	20.000000	79.000000	
50%	3.000000	117.000000	72.000000	23.000000	79.000000	
75%	6.000000	140.250000	80.000000	32.000000	127.250000	
max	17.000000	199.000000	122.000000	99.000000	846.000000	

	BMI	DiabetesPedigreeFunction	Age	Outcome
count	768.000000	768.000000	768.000000	768.000000
mean	32.450805	0.471876	33.240885	0.348958
std	6.875374	0.331329	11.760232	0.476951
min	18.200000	0.078000	21.000000	0.000000
25%	27.500000	0.243750	24.000000	0.000000
50%	32.000000	0.372500	29.000000	0.000000
75%	36.600000	0.626250	41.000000	1.000000
max	67.100000	2.420000	81.000000	1.000000

[196]:
```
# Countis the number of BloodPressure entries
# mean is the average of BloodPressure values
# std is the standard deviation (deviation from the mean) of BloodPressure␣
 ↪values
#min is the mininum value of BloodPressure values
# 25% is the 25 percentile mark
#50% is the 50 percentile mark(median) of BloodPressure values
# 75% is the 75 percentilemark of BloodPressurevalues
# max is the maximum value of BloodPressure values
diabetes["BloodPressure"].describe()
```

[196]:
```
count    768.000000
mean      72.250000
std       12.117203
min       24.000000
```

```
25%       64.000000
50%       72.000000
75%       80.000000
max      122.000000
Name: BloodPressure, dtype: float64
```

[198]:
```
diabetes.BloodPressure
```

[198]:
```
0      72
1      66
2      64
3      66
4      40
       ..
763    76
764    70
765    72
766    60
767    70
Name: BloodPressure, Length: 768, dtype: int64
```

[214]:
```
# Estimate the mean and standard deviation of BloodPressure values
mu = diabetes["BloodPressure"].mean()
print("The estimated mean is", round(mu,2))
```

```
The estimated mean is 72.25
```

[215]:
```
sigma = diabetes["BloodPressure"].std()
print("The estimated standard deviation is", round(sigma, 2))
```

```
The estimated standard deviation is 12.12
```

[210]:
```
# Q 10. Plot the distribution plot for the variable 'BloodPressure'. Write␣
 ↪detailed observations from the plot.

# Below, as we can see, theblue curve display the shape of sata distribution␣
 ↪while the red curve the PDF
# (Probabilitydensity function). It is obvious, the dataset is approsimately␣
 ↪normal. Therefore, our assumption is that
# the data distribution from the dataset to be normal and normality assumption␣
 ↪is what our calculations will be based upon.

density = pd.DataFrame()
density["x"] = np.linspace(diabetes["BloodPressure"].min() - 0.01,␣
 ↪diabetes["BloodPressure"].max() + 0.01, 100)
density["pdf"]= norm.pdf(density["x"], mu, sigma)

fig, ax = plt.subplots()
```

```
# plot the distribution of data using histogram
sns.histplot(diabetes["BloodPressure"], ax=ax, kde=True, stat="density")

# plot the pdf of the normal distribution
ax.plot(density["x"], density["pdf"], color="red")
plt.title("Normal Distribution")
plt.show()
```

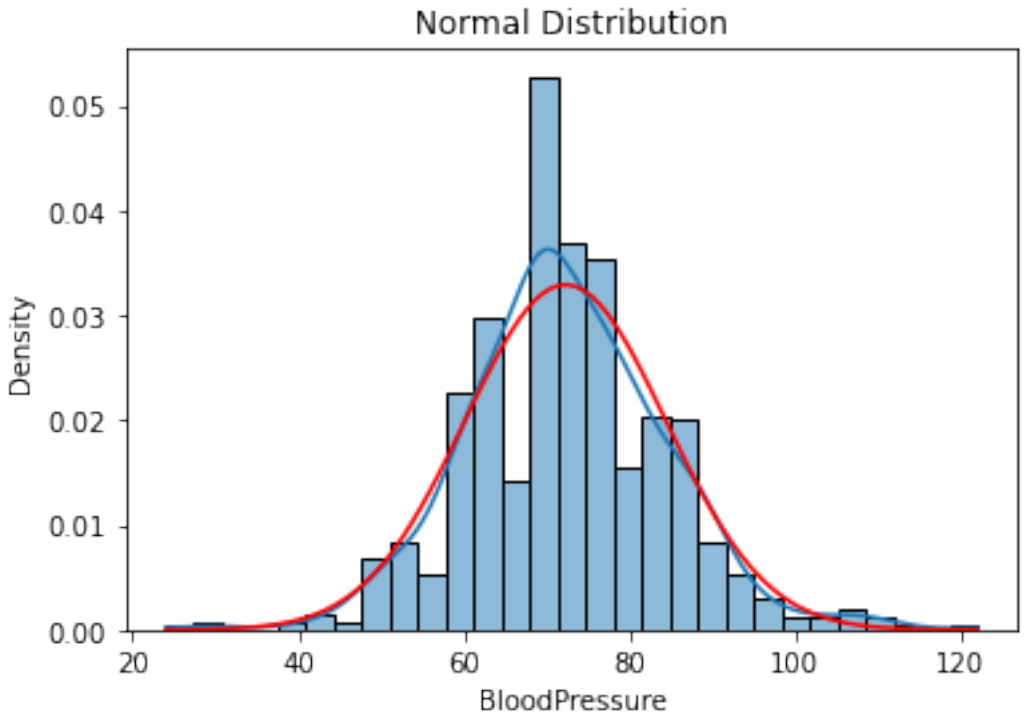

[218]:
```
# Q 11. What is the 'BMI' for the person having the highest 'Glucose'?

# The BMI for the person having the highest Glucose is 67.10

diabetes["BMI"].describe()
```

[218]:
```
count    768.000000
mean      32.450805
std        6.875374
min       18.200000
25%       27.500000
50%       32.000000
75%       36.600000
max       67.100000
Name: BMI, dtype: float64
```

[226]:
```
# Q 12. Q 12.1 What is the mean of the variable 'BMI'?
#12.3 What is the mean of the variable 'BMI'?
diabetes["BMI"].mean()

#12.2 What is the median of the variable 'BMI'?
#median = 32
diabetes["BMI"].median()

#12.3 What is the mode of the variable 'BMI'?
diabetes["BMI"].mode()


#12.4 Are the three measures of central tendency equal?
# Yes, all three measures of central tendency are equal.
#mean = median = mode = 32.0
```

[226]:
```
0    32.0
Name: BMI, dtype: float64
```

[237]:
```
# use describe method to find the mean
diabetes["Glucose"].describe()
```

[237]:
```
count    768.000000
mean     121.675781
std       30.436252
min       44.000000
25%       99.750000
50%      117.000000
75%      140.250000
max      199.000000
Name: Glucose, dtype: float64
```

[245]:
```
# Sort the Glucose values in ascending orders

diabetes.sort_values('Glucose')
```

[245]:
```
     Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin   BMI  \
62             5       44             62             20       79  25.0
680            2       56             56             28       45  24.2
146            9       57             80             37       79  32.8
537            0       57             60             20       79  21.7
352            3       61             82             28       79  34.4
..           …       …              …              …       …   …
579            2      197             70             99       79  34.7
408            8      197             74             20       79  25.9
8              2      197             70             45      543  30.5
561            0      198             66             32      274  41.3
```

```
661            1      199             76             43       79  42.9

     DiabetesPedigreeFunction  Age  Outcome  Glucose_ranked
62                      0.587   36        0           768.0
680                     0.332   22        0           767.0
146                     0.096   41        0           766.0
537                     0.735   67        0           766.0
352                     0.243   46        0           764.0
..                        …  …       …              …
579                     0.575   62        1             6.0
408                     1.191   39        1             6.0
8                       0.158   53        1             6.0
561                     0.502   28        1             2.0
661                     1.394   22        1             1.0

[768 rows x 10 columns]
```

[246]:
```
# Q 13. How many women's 'Glucose' level is above the mean level of 'Glucose'?


# locate the mean = 121.675781 and rank it.

# The mean is located below "Glucose_ranked" of 343. So, there are 343 values␣
 ↪above the mean

diabetes.iloc[340:348]
```

[246]:
```
     Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin   BMI  \
340            1      130             70             13      105  25.9
341            1       95             74             21       73  25.9
342            1      120             68             35       79  32.0
343            5      122             86             20       79  34.7
344            8       95             72             20       79  36.8
345            8      126             88             36      108  38.5
346            1      139             46             19       83  28.7
347            3      116             69             20       79  23.5

     DiabetesPedigreeFunction  Age  Outcome  Glucose_ranked
340                     0.472   22        0           258.0
341                     0.673   36        0           626.0
342                     0.389   22        0           365.0
343                     0.290   33        0           343.0
344                     0.485   57        0           626.0
345                     0.349   49        0           297.0
346                     0.654   22        0           205.0
347                     0.187   23        0           400.0
```

[257]:
```
#Q 14. How many entries (women) have their 'BloodPressure' equal to the median
 ↪of 'BloodPressure'
# and their 'BMI' less than the median of 'BMI'?

# There are 113 entries (women) that their 'BloodPressure' equal to the median
 ↪of 'BloodPressure'
# and their 'BMI' less than the median of 'BMI'


# BloodPressure median = 72.0

# 1st, sort BloodPressure
# Second, locate all values of median =72, then sum the corresponding women's
 ↪entries that are equal to the BloodPlressure
# of median 72.0.

# So, there are 234 entries (women)
diabetes["BloodPressure"].median()
```

[257]:
```
72.0
```

[255]:
```
diabetes["BMI"].median()
```

[255]:
```
32.0
```

[285]:
```
ranked = diabetes.sort_values(['BloodPressure','BMI'])
```

[286]:
```
ranked
```

[286]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
597	1	89	24	19	25	27.8	
18	1	103	30	38	83	43.3	
125	1	88	30	42	99	55.0	
599	1	109	38	18	120	23.1	
4	0	137	40	35	168	43.1	
..	…	…	…	…	…	…	
549	4	189	110	31	79	28.5	
43	9	171	110	24	240	45.4	
177	0	129	110	46	130	67.1	
691	13	158	114	20	79	42.3	
106	1	96	122	20	79	22.4	

	DiabetesPedigreeFunction	Age	Outcome	Glucose_ranked
597	0.559	21	0	675.0
18	0.183	33	0	537.0
125	0.496	26	1	684.0
599	0.407	26	0	471.0

```
4                        2.288   33        1           218.0
..                         …   …       …             …
549                      0.680   37        0            22.0
43                       0.721   54        1            69.0
177                      0.319   26        1           272.0
691                      0.257   44        1           112.0
106                      0.207   27        0           613.0

[768 rows x 10 columns]
```

[287]:
```
ranked.iloc[19:349]
```

[287]:
```
     Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin   BMI  \
258            1      193             50             16      375  25.9
243            6      119             50             22      176  27.1
687            1      107             50             19       79  28.3
98             6       93             50             30       64  28.7
313            3      113             50             10       85  29.5
..           …      …             …              …       …   …
515            3      163             70             18      105  31.6
262            4       95             70             32       79  32.1
570            3       78             70             20       79  32.5
191            9      123             70             44       94  33.1
241            4       91             70             32       88  33.1

     DiabetesPedigreeFunction  Age  Outcome  Glucose_ranked
258                     0.655   24        0            16.0
243                     1.318   33        1           376.0
687                     0.181   29        0           495.0
98                      0.356   23        0           640.0
313                     0.626   25        0           426.0
..                        …   …       …             …
515                     0.268   28        1            92.0
262                     0.612   24        0           626.0
570                     0.270   39        0           739.0
191                     0.374   40        0           331.0
241                     0.446   22        0           658.0

[330 rows x 10 columns]
```

[289]:
```
# Q 15. Below is the pairplot of variables 'Glucose', 'SkinThickness' and␣
 ↪'DiabetesPedigreeFunction'.
#Write you observations from the plot.

# My observation is that the variables take many different forms of␣
 ↪distribution. It shows the relationship
# with each other.Some show histogram while others are skewed.
```

```
diabetes.head(20)
```

[289]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	6	148	72	35	79	33.600000	
1	1	85	66	29	79	26.600000	
2	8	183	64	20	79	23.300000	
3	1	89	66	23	94	28.100000	
4	0	137	40	35	168	43.100000	
5	5	116	74	20	79	25.600000	
6	3	78	50	32	88	31.000000	
7	10	115	69	20	79	35.300000	
8	2	197	70	45	543	30.500000	
9	8	125	96	20	79	31.992578	
10	4	110	92	20	79	37.600000	
11	10	168	74	20	79	38.000000	
12	10	139	80	20	79	27.100000	
13	1	189	60	23	846	30.100000	
14	5	166	72	19	175	25.800000	
15	7	100	69	20	79	30.000000	
16	0	118	84	47	230	45.800000	
17	7	107	74	20	79	29.600000	
18	1	103	30	38	83	43.300000	
19	1	115	70	30	96	34.600000	

	DiabetesPedigreeFunction	Age	Outcome	Glucose_ranked
0	0.627	50	1	148.0
1	0.351	31	0	701.0
2	0.672	32	1	35.0
3	0.167	21	0	675.0
4	2.288	33	1	218.0
5	0.201	30	0	400.0
6	0.248	26	1	739.0
7	0.134	29	0	410.0
8	0.158	53	1	6.0
9	0.232	54	1	311.0
10	0.191	30	0	459.0
11	0.537	34	1	76.0
12	1.441	57	0	205.0
13	0.398	59	1	22.0
14	0.587	51	1	82.0
15	0.484	32	1	576.0
16	0.551	31	1	382.0
17	0.254	31	1	495.0
18	0.183	33	0	537.0
19	0.529	32	1	410.0

[290]:
```
#set the figure size
plt.figure(figsize = (11,11))
```

[290]: <Figure size 792x792 with 0 Axes>

<Figure size 792x792 with 0 Axes>

[291]:
```
#plot a pairt plot
sns.pairplot(diabetes)
```

[291]: <seaborn.axisgrid.PairGrid at 0x291c136b4f0>

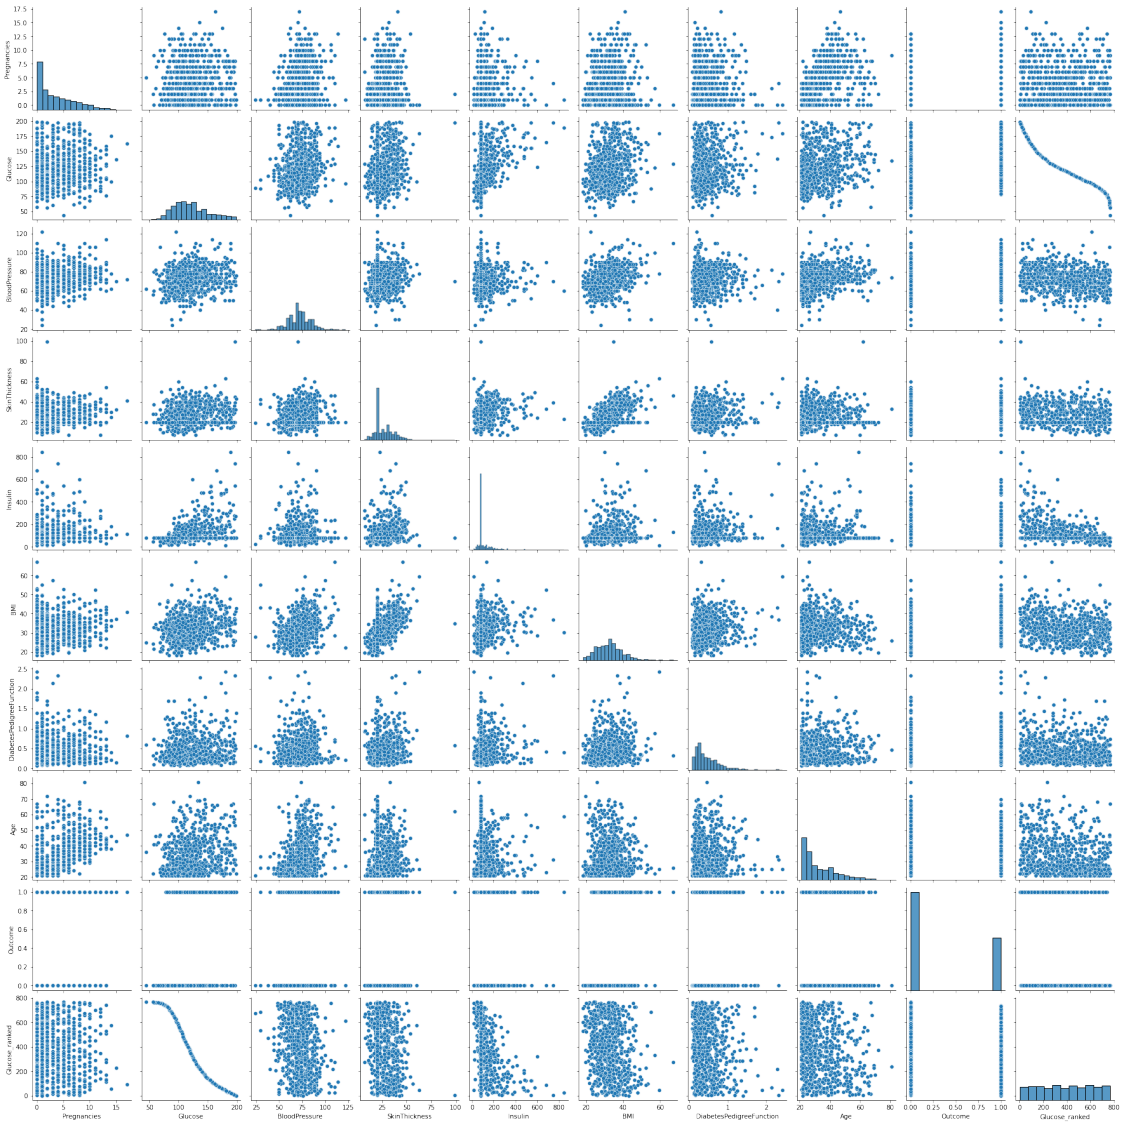

[293]:
```
plt.show()
```

[294]:
```
# Q 16. Plot the scatterplot between 'Glucose' and 'Insulin'. Write your␣
 ↪observations from the plot.

# The idea from the scatter plot,the more glucose,the more insulin. There is a␣
 ↪rough positive correlation between the two

#data
X = diabetes['Glucose']
Y = diabetes['Insulin']

#Plot the scatter plot
plt.scatter(X,Y)

# add the axes labels to the plot
plt.xlabel('Glucose')
plt.ylabel('Insulin')

# display the plot
plt.show()
```

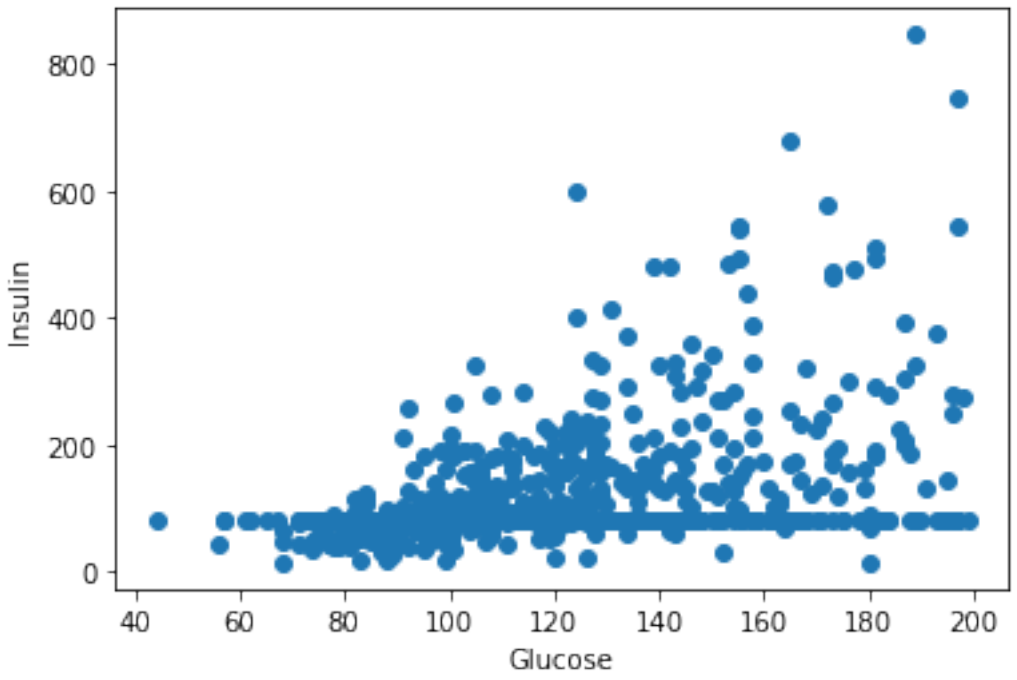

[296]:
```
# Q 17. Plot the boxplot for the 'Age' variable. Are there outliers?

Yes, they are outliers: these points outside the whisker
```

```
# plot a distribution of Age
plt.boxplot(diabetes['Age'])

# add labels for five numbersummary
plt.text(x = 1.1, y = diabetes['Age'].min(), s='min')
plt.text(x = 1.1, y = diabetes.Age.quantile(0.25), s ='Q1')
plt.text(x = 1.1, y = diabetes['Age'].median(), s ='meadian (Q2)')
plt.text(x = 1.1, y = diabetes.Age.quantile(0.75), s ='Q3')
plt.text(x = 1.1, y = diabetes['Age'].max(), s ='max')

# ass the graphtitle andaxes labels
plt.title('Boxplot of Age')
plt.ylabel('Age')

# display the plot
plt.show()
```

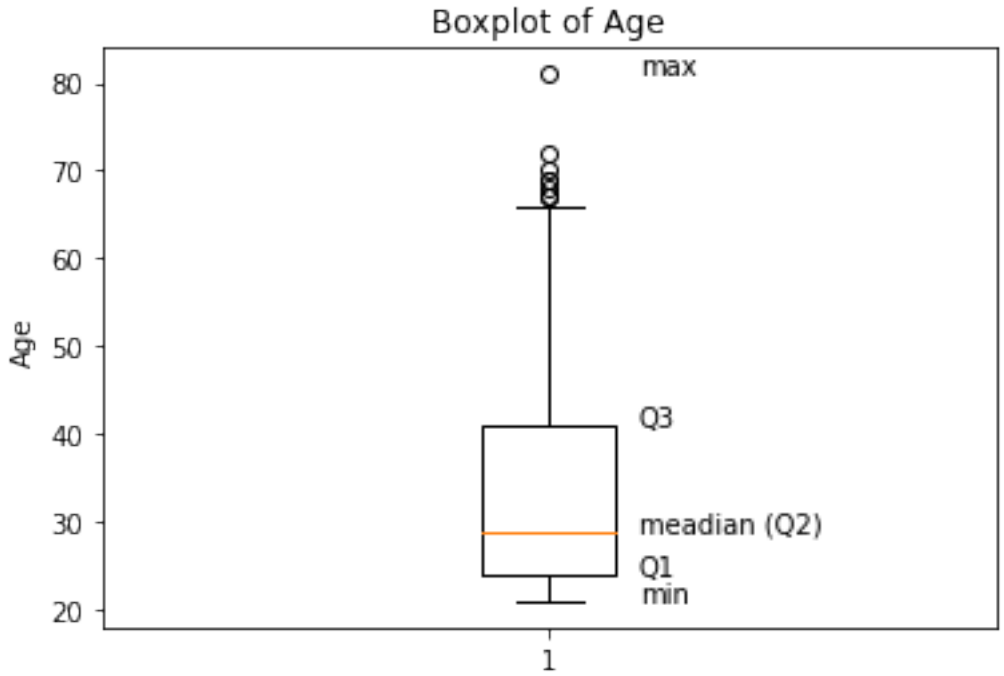

[298]:
```
# Q 18. Plot histograms for variable Age to understand the number of women in␣
↪different Age groups given that they have diabetes or not.
# Explain both histograms and compare them.

# The Age variable is positively skewed.
```

```
# plot the histogram
# specify the number of bins, using 'bins' parameter
plt.hist(diabetes['Age'], bins = 5)

# add the graph title and axes labels
plt.title('Distribution of Age')
plt.xlabel('Age')
plt.ylabel('Frequency')

# display theplot
plt.show()
```

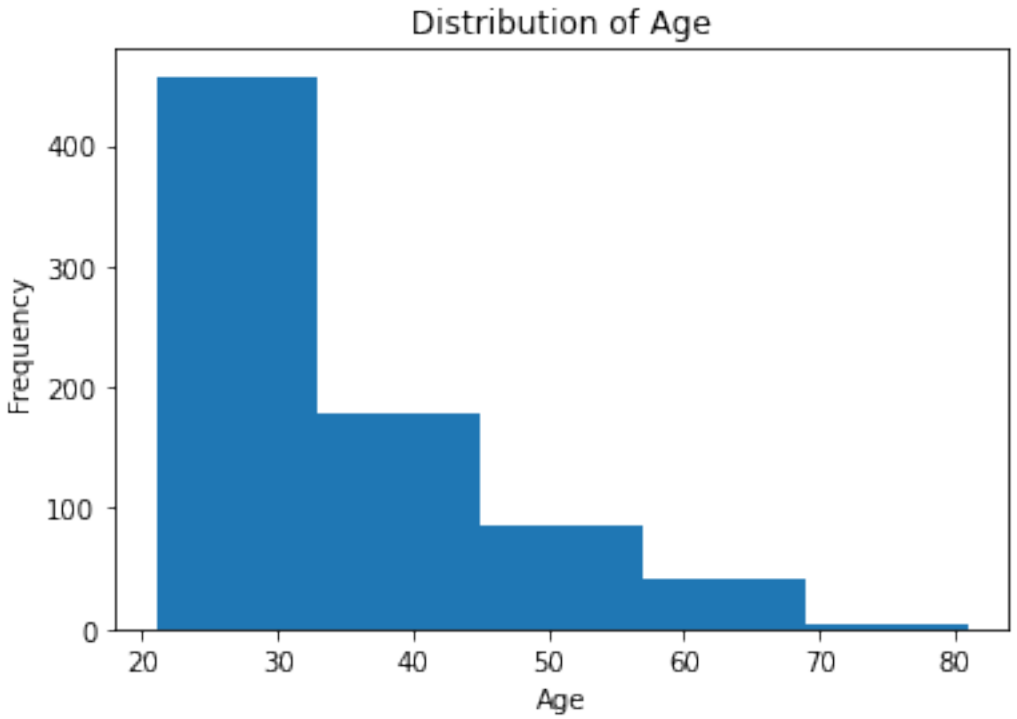

[299]:
```
# Q 19. What is Inter Quartile Range of all the variables? Why is it used? Which plot visualizes the same?
# Interquatile IQR = Q3 - Q1 = 41.0000 - 24.000 = 17.000.

# Interquatile Range is used, in a modified boxplot, to represent outliers as special points. The value of IQR can be used
# to idetify outliers as follows: above Q3 by anamount greater than 1.5 X IQR or below Q1 by an amount greater than 1.5 X IQR

diabetes["Age"].describe()
```

[299]:
```
count    768.000000
mean      33.240885
std       11.760232
min       21.000000
25%       24.000000
50%       29.000000
75%       41.000000
max       81.000000
Name: Age, dtype: float64
```

[300]:
```
# Q 20. Find and visualize the the correlation matrix. Write your observations␣
 ↪from the plot.

# The corrrelation matrix shows correlation coefficients between sets of␣
 ↪variables in the dataset. For instance some
# varables could be postively correlated while others could be negatively␣
 ↪correlated. Pregancies and insulin are negative
# correlated.
corrM = diabetes.corr()
```

[301]:
```
corrM
```

[301]:

	Pregnancies	Glucose	BloodPressure	SkinThickness
Pregnancies	1.000000	0.128022	0.208987	0.009393
Glucose	0.128022	1.000000	0.219765	0.158060
BloodPressure	0.208987	0.219765	1.000000	0.130403
SkinThickness	0.009393	0.158060	0.130403	1.000000
Insulin	-0.018780	0.396137	0.010492	0.245410
BMI	0.021546	0.231464	0.281222	0.532552
DiabetesPedigreeFunction	-0.033523	0.137158	0.000471	0.157196
Age	0.544341	0.266673	0.326791	0.020582
Outcome	0.221898	0.492884	0.162879	0.171857
Glucose_ranked	-0.137137	-0.973619	-0.235622	-0.151535

	Insulin	BMI	DiabetesPedigreeFunction
Pregnancies	-0.018780	0.021546	-0.033523
Glucose	0.396137	0.231464	0.137158
BloodPressure	0.010492	0.281222	0.000471
SkinThickness	0.245410	0.532552	0.157196
Insulin	1.000000	0.189919	0.158243
BMI	0.189919	1.000000	0.153508
DiabetesPedigreeFunction	0.158243	0.153508	1.000000
Age	0.037676	0.025748	0.033561
Outcome	0.178696	0.312254	0.173844
Glucose_ranked	-0.386221	-0.232370	-0.119526

	Age	Outcome	Glucose_ranked

```
Pregnancies                0.544341  0.221898       -0.137137
Glucose                    0.266673  0.492884       -0.973619
BloodPressure              0.326791  0.162879       -0.235622
SkinThickness              0.020582  0.171857       -0.151535
Insulin                    0.037676  0.178696       -0.386221
BMI                        0.025748  0.312254       -0.232370
DiabetesPedigreeFunction   0.033561  0.173844       -0.119526
Age                        1.000000  0.238356       -0.274993
Outcome                    0.238356  1.000000       -0.481950
Glucose_ranked            -0.274993 -0.481950        1.000000
```

[ ]: